

What is Artificial Intelligence?

- Artificial Intelligence is the synthesis and analysis of computational agents that act intelligently.
- An agent is something that acts in an environment.
- An agent acts intelligently if:
 - ▶ its actions are appropriate for its goals and circumstances; or
 - ▶ it is flexible to changing; or environments and goals; or
 - ▶ it learns from experience; or
 - ▶ it makes appropriate choices given perceptual and computational limitations.


Examples of Agents

- **Animals**: human, dog, mouse, bird, insect, worm, ...
- **Biological cells**: bacterium, neuron, ...
- **Organisations**: Microsoft, NZ Government, CSSE Department, UC, ...
- **Software and devices**: thermostat, user interface, airplane controller, network controller, game, advising system, tutoring system, diagnostic assistant, robot, Google car, Mars rover, ...

Goals of Artificial Intelligence

- **Scientific goal:** to understand the principles that make intelligent behavior possible in natural or artificial systems.
 - ▶ analyze natural and artificial agents;
 - ▶ formulate and test hypotheses about what it takes to construct intelligent agents;
 - ▶ design, build, and experiment with computational systems that perform tasks that require intelligence.
- **Engineering goal:** design useful, intelligent artifacts.

Agents acting in an environment




Inputs to an agent

- **Abilities** — the set of possible actions it can perform
- **Goals/Preferences** — what it wants, its desires, its values,...
- **Prior Knowledge** — what it comes into being knowing, what it doesn't get from experience,...
- **History** of stimuli
 - ▶ (current) **stimuli** — what it receives from environment now (observations, percepts)
 - ▶ **past experiences** — what it has received in the past

Example agent: autonomous car

- **abilities:** steer, accelerate, brake
- **goals:** safety, get to destination, timeliness,...
- **prior knowledge:** what signs mean, what to stop for
- **stimuli:** vision, laser, GPS...
- **past experiences:** streetmaps, how breaking, steering affects direction..

Domain for Delivery Robot




Autonomous Delivery Robot

Example inputs:

- **Abilities:** movement, speech, pickup and place objects.
- **Prior knowledge:** its capabilities, objects it may encounter, maps.
- **Past experience:** which actions are useful and when, what objects are there, how its actions affect its position.
- **Goals:** what it needs to deliver and when, tradeoffs between acting quickly and acting safely.
- **Stimuli:** about its environment from cameras, sonar, sound, laser range finders, or keyboards.

Domain for Diagnostic Assistant in a Smart Home



Diagnostic Assistant

Example inputs:


- **Abilities:** recommends fixes, ask questions.
- **Prior knowledge:** how switches and lights work, how malfunctions manifest themselves, what information tests provide, the side effects of repairs.
- **Past experience:** the effects of repairs or treatments, the prevalence of faults or diseases.
- **Goals:** fixing the device and tradeoffs between fixing or replacing different components.
- **Stimuli:** symptoms of a device or patient.

Common Tasks of the Domains

- **Modeling the environment** Build models of the physical environment, patient, or information environment.
- **Evidential reasoning or perception** Given observations, determine what the world is like.
- **Action** Given a model of the world and a goal, determine what should be done.
- **Learning from past experiences** Learn about the specific case and the population of cases.

Agent System Architecture

An **agent** is made up of a **body** and a **controller**.




- An agent interacts with the environment through its body.
- The **body** is made up of:
 - ▶ **sensors** that interpret stimuli
 - ▶ **actuators** that carry out actions
- The controller receives **percepts** from the body.
- The controller sends **commands** to the body.
- The body can also have reactions that are not controlled.


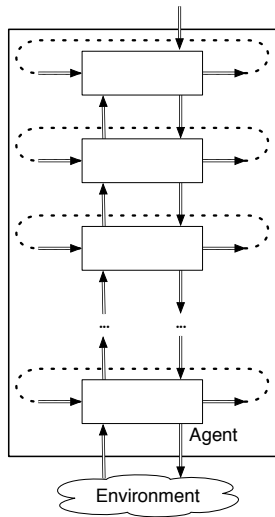
Implementing a controller

- A **controller** is the **brains** of the agent.
- Agents are situated in time, they receive sensory data in time, and do actions in time.
- Controllers have (limited) memory and (limited) computational capabilities.
- The controller specifies the command at every time.
- The command at any time can depend on the current and previous percepts.

Controller



Hierarchical Robotic System Architecture



Example: delivery robot

- The robot has three actions: go straight, go right, go left. (Its velocity doesn't change).
- It can be given a **plan** consisting of sequence of named locations for the robot to go to in turn.
- The robot must avoid obstacles.
- It has a single **whisker sensor** pointing forward and to the right. The robot can detect if the whisker hits an object. The robot knows where it is.
- The obstacles and locations can be moved dynamically. Obstacles and new locations can be created dynamically.

A Decomposition of the Delivery Robot

