

Local and Global Search

Parts adapted from:

- Chapter 4 of AI 2E by David Poole and Alan Macworth;
- AI a modern approach by Stuart Russel and Peter Norvig

Optimisation problems

Optimisation problem:

- A set of variables and their domains
- An objective function

Find an assignment that optimises (maximise / minimises) the value of the objective function.


A **constrained optimisation problem**, in addition to the above, has a set of constraints that determine what assignments are allowed. In a constrained optimisation problem, the goal is to find an assignment that satisfies the constraints and optimises the objective function.

Local search

- Optimisation usually involves searching.
- Like CSP, the path is irrelevant; only the solution matters.
- In local search we use algorithms that *iteratively improve* a state.
- We keep a single **current state**, and in each iteration we try to improve it by moving to one of its **neighbours**.
- This takes constant space.
- The goal is to find an optimal state.
- Most local search algorithms are greedy.
- Two common local search algorithms are **hill climbing** (greedy ascent for maximisation) and **greedy descent** (for minimisation).

Example: Traveling Salesperson Problem

- Start with any complete tour, and perform pairwise exchanges
- Variants of this approach get very close to optimal solution very quickly with large numbers of cities.




Local search for CSPs

- A constrained satisfaction problem (CSP) can be reduced to an optimisation problem.
- Aim is to find an assignment with zero unsatisfied constraints.
- Given an assignment of a value to each variable, a **conflict** is an unsatisfied constraint.
- The goal is an assignment with zero conflicts.
- Heuristic function to be minimised: the number of conflicts.

Local search: n -queens problem


- Aim: Put n queens on an $n \times n$ board with no two queens attacking each other.
- The objective (heuristic) function to minimise: number of conflicts.



$$h = 1$$

Example: 4-Queens

- States: 4 queens in 4 columns ($4^4 = 256$ states)
- Obtaining neighbours: move queen in column
- Objective function to minimise: $h(n)$ = number of pairs of queens that are attacking each other (number of conflicts)



Example: neighbours

- Objective function (conflict count): number of pairs of queens that are attacking each other.
- Number of conflicts in the current state: 17


18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♚	13	16	13	16
♚	14	17	15	♚	14	16	16
17	♚	16	18	15	♚	15	♚
18	14	♚	15	15	14	♚	16
14	14	13	17	12	14	12	18

Some Variants of Greedy Descent

- Find the variable-value pair that minimises the number of conflicts at every step.
- Select a variable that participates in the most number of conflicts. Select a value that minimises the number of conflicts.
- Select a variable that appears in any conflict. Select a value that minimises the number of conflicts.

Local Search Issues

- Local search can get stuck in local optima or flat areas of the **landscape** of the objective function.
- Randomised greedy descent can help sometimes:
 - **random step**: move to a random neighbour.
 - **random restart**: reassign random values to all variables.
- these make the search global.



a local minimum with a conflict count of 1.

Parallel search

- A total assignment is called an **individual**.
- Idea: maintain a population of k individuals instead of one.
- At every stage, update each individual in the population. Whenever an individual is a solution, it can be reported.
- Like k restarts, but uses k times the minimum number of steps.
- A basic form of global search.

Simulated Annealing

- Pick a variable at random and a new value at random.
- If it is an improvement, adopt it.
- If it isn't an improvement, adopt it probabilistically depending on a temperature parameter, T .
 - ▶ With current assignment n and proposed assignment n' we move to n' with probability $e^{(h(n)-h(n'))/T}$
- Temperature can be reduced.


Probability of accepting a change:

Temperature	1-worse	2-worse	3-worse
10	0.91	0.81	0.74
1	0.37	0.14	0.05
0.25	0.02	0.0003	0.000005
0.1	0.00005	0	0

Gradient Descent

- A widely-used local search algorithm in numeric optimisation (e.g. in machine learning)
- Used when the variables are numeric and continuous.
- The objective function must be differentiable (mostly).

```
1: Guess  $\mathbf{x}^{(0)}$ , set  $k \leftarrow 0$   
2: while  $\|\nabla f(\mathbf{x}^{(k)})\| \geq \epsilon$  do  
3:    $\mathbf{x}^{(k+1)} = \mathbf{x}^{(k)} - t_k \nabla f(\mathbf{x}^{(k)})$   
4:    $k \leftarrow k + 1$   
5: end while  
6: return  $\mathbf{x}^{(k)}$ 
```



Evolutionary Algorithms

References:

A.E. Eiben and J.E. Smith,
Introduction to Evolutionary
Computing, Springer

K. A. De Jong, Evolutionary
Computation, MIT Press

J. C. Spall
*Introduction to Stochastic Search
and Optimization*, John Wiley and
Sons



Genetic (Evolutionary) Algorithms





- Inspired by natural selection
- A form of global search
- Requires:
 - Representation
 - Evaluation function
 - Selection of parents
 - Reproduction operators
 - Initialisation procedure
 - Parameter settings

Evaluation: Fitness Function


Purpose:

- Parent selection
- Measure for convergence
- For Steady state: Selection of individuals to die
- Should reflect the value of the chromosome in some “real” way
- It is a critical part of any EA / GA

Fitness landscapes



GA: Flowchart



Selection


Main idea: better individuals should have higher chance of surviving and breeding.

Types:

- Roulette wheel selection
- Tournament selection
- ... any mechanism that somehow overall achieves the main idea.

Roulette Wheel Selection

- Chances proportional to fitness
- Assign to each individual a part of the roulette wheel
- Spin the wheel n times to select n individuals



Individual	Fitness
1	1.0
2	2.0
3	3.0
4	4.0
5	5.0

Roulette Wheel Selection: Example

- Sum the fitness of all individuals, call it T
- Generate a random number N between 1 and T
- Return individual whose fitness added to the running total is equal to or larger than N
- Chance to be selected is exactly proportional to fitness
- Individual: 1, 2, 3, 4, 5, 6
- Fitness: 8, 2, 17, 7, 4, 11
- Running total: 8, 10, 27, 34, 38, 49
- N: 23
- Selected: 3

Selection: Tournaments

- n individuals are randomly chosen; the fittest one is selected as a parent.
- n is the “size” of the tournament.
- By changing the size, selection pressure can be adjusted.

Elitism

- Widely used in population models
- Always keep at least one copy of the fittest solution so far
- Results in non-decreasing (maximum) fitness over generation

Reproduction Operators

Crossover

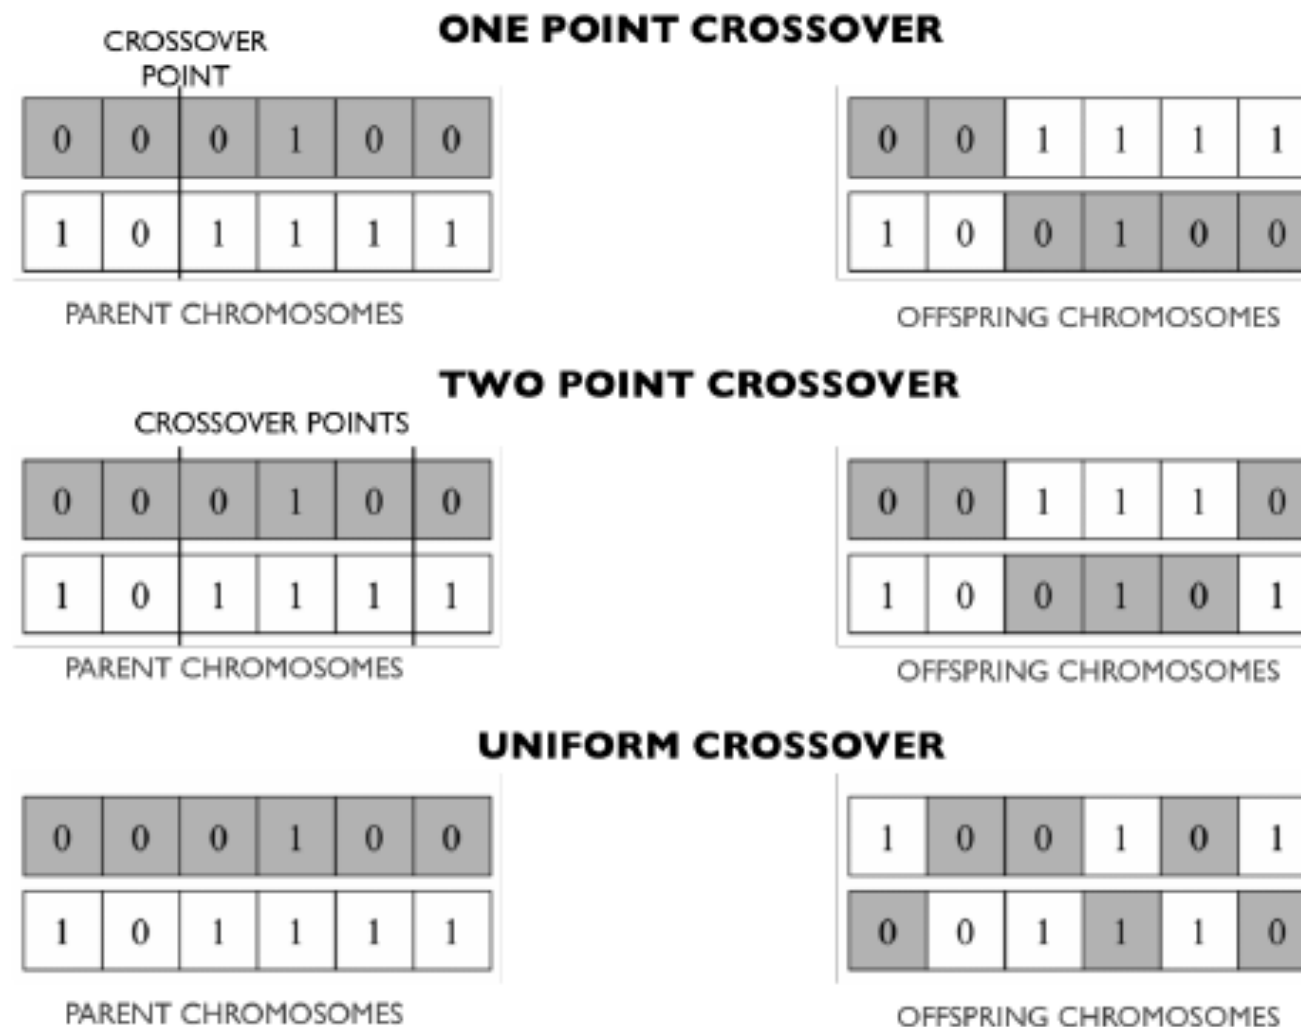
- Two parents produce two offspring
- There is a chance that the chromosomes of the two parents are copied unmodified as offspring
- There is a chance that the chromosomes of the two parents are randomly recombined (crossover) to form offspring
- Typically the chance of crossover is between 0.6 and 1.0
- Types: 1-point, 2-point, Uniform, ...

Mutation

- There is a chance that a gene of a child is changed randomly
- Typically the chance of mutation is low (e.g. 0.001)

Crossover


- Generate 1, 2, or a number of random crossover points.
- Split the parents at these points.
- Create offsprings by exchanging alternate segments.



Crossover or Mutation?


- Purpose of crossover: combining somewhat good candidates in the hope of producing better children
- Purpose of mutation: bring diversity (new ideas!)
- Decade long debate: which one is better/necessary?
- A rather wide agreement:
 - it depends on the problem, but
 - in general, it is good to have both.
 - mutation-only-EA is possible, crossover-only-EA would not work.

Tree representation



```


i = 1;
while (i < 20)
{
    i = i + 1
}
  
```



$$2 \cdot \pi + \left((x + 3) - \frac{y}{5 + 1} \right)$$


Randomly Generating Programs

(+ ...)




Randomly Generating Programs

(+ 2 ...)




Randomly Generating Programs

(+ 2 3 ...)




Randomly Generating Programs

(+ 2 3 (* ...) ...)




Randomly Generating Programs

(+ 2 3 (* X 7) (/ ...))




Randomly Generating Programs

(+ 2 3 (* X 7) (/ Y 5))



Mutation


(+ 2 3 (* X 7) (/ Y 5))



Mutation

(+ 2 3 (* X 7) (/ Y 5))


First pick a random node



Mutation


(+ 2 3 (+ (* 4 2) 3) (/ Y 5))

Delete the node and its children, and replace with a randomly generated program




Crossover

$(+ X (* 3 Y))$




$(- (/ 25 X) 7)$




Crossover

(+ X (* 3 Y))




Pick a random node in each program

(- (/ 25 X) 7)



Crossover

(+ X (* (/ 25 X) Y))



Swap the two nodes

(- 3 7)

